

[Pages](#) / ... / [Tools Tips and Tricks](#)

MPLAB XC16 and XC-DSC DWARF Differences

Created by Calum Wilkie, last modified less than a minute ago

Although XCxx is based upon GCC and uses the common GCC DWARF generation code, the ELF container that holds this information has some oddities. These oddities are mostly related to the way we store data in the ELF file. Here we attempt to highlight these differences so that our DWARF files may be properly decoded.

In this tutorial I will be using the XCxx compiler and some of the binary utilities included with the compiler to inspect resulting object files. I will also be using the unix tool 'od' (octal dump) to dump raw files; any other appropriate tool (for Windows users) can be substituted.

XC16 vs XC-DSC

In this document I will refer to XC16 and XC-DSC together or interchangeably. XC-DSC is the successor to XC16 targeting Microchip DSC devices. Any command line that mentions xc16 or xc-dsc, can be substituted with the other command line.

dsPIC30F and similar devices

ELF Sections

Most ELF sections are padded with null bytes. This is usually thought of for 'Flash' type sections, so that we can fit a 24-bit word into a 16-bit address hole. An oddity of the architecture is that the address increments by 2 for each program word, and each program word is 24-bits wide. Therefore a word in program memory is represented with 4 bytes of data. For example, the program word 0x112233 is encoded as 0x112233xx, where xx (the padding byte) is typically null, although a customer might specify a different padding byte value with a command option.

To allow for address compatibility between data memory and program memory, each word in data memory (16 bits) is also represented with 4 bytes of data. This is accomplished by inserting a padding byte between each original data byte. For example, the data word 0x1122 is encoded as 0x11002200. All data sections are represented in this way, including DWARF debugging sections.

If an external tool is used to interpret DWARF debugging information for PIC24, dsPIC30F, or dsPIC33C/E/F devices, then the padding bytes must be recognized and discarded.

A Short Example

Consider this short example (dwarf.c):

```
int an_int = 0xF00D;

char *a_string = "This is a string";

main() {}
```

In this example, an_int and a_string are data objects, while the value pointed to by a_string is usually located in Flash. main() exists solely to get a linked executable.

Compiling this file (with -save-temps) will leave behind dwarf.o and a.out (and some other detritus); be sure to use -g to get relevant DWARF information. We can use xc16-objdump to inspect both the object file and executable file (these are both the same kind of ELF object file). The tool xc16-readelf can be used to inspect these ELF files in a more specific way and also to decode the DWARF information into a more readable form. For our purposes today, this will not be as helpful.

I compiled this file with: `xc16-gcc -g -save-temps dwarf.c` to produce `dwarf.o` and `a.out` I then use `xc16-objdump -h` to display the section information for the object file:

```
dwarf.o:      file format elf32-pic30
Sections:
```

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	00000006	00000000	00000000	00000034	2**1
	CONTENTS, ALLOC, LOAD, CODE					
1	.data	00000000	00000000	00000000	00000040	2**1
	ALLOC, LOAD, DATA					
2	.bss	00000000	00000000	00000000	00000040	2**1
	ALLOC					
3	.debug_abbrev	0000004e	00000000	00000000	00000040	2**0
	CONTENTS, DEBUGGING					
4	.debug_info	000000bb	00000000	00000000	000000dc	2**0
	CONTENTS, RELOC, DEBUGGING					
5	.debug_line	00000033	00000000	00000000	00000252	2**0
	CONTENTS, RELOC, DEBUGGING					
6	.ndata	00000004	00000000	00000000	000002b8	2**1
	CONTENTS, ALLOC, LOAD, RELOC, DATA, NEAR					
7	.const	00000012	00000000	00000000	000002c0	2**1
	CONTENTS, ALLOC, LOAD, READONLY, PSV, PAGE					
8	.debug_frame	0000002e	00000000	00000000	000002e4	2**0
	CONTENTS, RELOC, DEBUGGING					
9	.debug_pubnames	00000033	00000000	00000000	00000340	2**0
	CONTENTS, RELOC, DEBUGGING					
10	.debug_aranges	00000018	00000000	00000000	000003a6	2**0
	CONTENTS, RELOC, DEBUGGING					
11	.debug_str	00000000	00000000	00000000	000003d6	2**0
	DEBUGGING					
12	__c30_info	00000000	00000000	00000000	000003d6	2**0
	DEBUGGING					
13	__c30_signature	00000006	00000000	00000000	000003d6	2**0
	CONTENTS, DEBUGGING					

We can use the options `-s` and `-j <name>` to dump the raw data for a particular section. I will dump the data for section `.ndata` (a near data section which should contain an `_int`) and `.const` (the flash const section which should contain the data for our string):

```
dwarf.o:      file format elf32-pic30
Contents of section .ndata:
 0000 0df00000          ....
```

```
dwarf.o:      file format elf32-pic30
Contents of section .const:
 0000 54686973 20697320 61207374 72696e67  This is a string
 0010 0000                ..
```

I have high-lit the value for an `_int`; the string is quite visible.

The 'File off' output from `xc16-objdump` gives the raw offset for the data for each file. Using a hex editor, or `od` (octal dump), we can inspect the raw data at each offset. Here is the offset for the `.ndata` section; the addresses are in hex and line is displayed in single byte hex and single byte (escaped) ASCII, I have highlighted the raw data starting at `0x2B8` (there are 4 bytes):

```
00002b0  04 00 00 00 01 00 01 00 0d 00 f0 00 00 00 00 00
          004 \0 \0 \0 001 \0 001 \0 \r \0 360 \0 \0 \0 \0 \0
```

And for the `.const` section:

```
00002c0  54 00 68 00 69 00 73 00 20 00 69 00 73 00 20 00
          T \0 h \0 i \0 s \0 \0 i \0 s \0 \0
00002d0  61 00 20 00 73 00 74 00 72 00 69 00 6e 00 67 00
          a \0 \0 s \0 t \0 r \0 i \0 n \0 g \0
00002e0  00 00 00 00 10 00 00 00 00 00 00 00 ff 00 ff 00
          \0 \0 \0 \0 020 \0 \0 \0 \0 \0 \0 \0 \0 377 \0 377 \0
```

In each case, the real data is followed by a byte of padding (in green).

The content for a.out (the linked executable) is similar, but addresses have been filled in.

DWARF Content

How does this relate to the DWARF content? DWARF information is content, like any other, and is stored as ELF sections in a similar way.

debug_abbrev:

Contains padding bytes:

```
0000040  01 00 11 00 01 00 25 00 08 00 13 00 0b 00 03 00
          001 \0 021 \0 001 \0 % \0 \b \0 023 \0 \v \0 003 \0
```

debug_info:

Hard to see, but ... contains padding bytes:

```
00000d0  0b 00 49 00 13 00 00 00 00 00 00 00 00 b7 00 00 00
          \v \0 I \0 023 \0 \0 \0 \0 \0 \0 \0 267 \0 \0 \0
00000e0  00 00 00 00 02 00 00 00 00 00 00 00 00 00 00 00 00
          \0 \0 \0 \0 002 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0
00000f0  04 00 01 00 47 00 4e 00 55 00 20 00 43 00 20 00 00
          004 \0 001 \0 G \0 N \0 U \0 \0 \0 c \0 \0
```

and so on...

Comparing Raw data with xc16-readelf

Using xc16-readelf to dump the .debug_info section into a readable form, I will highlight one particular abbreviation:

```
<1><62>: Abbrev Number: 2 (DW_TAG_subprogram)
  DW_AT_external      : 1
  DW_AT_name          : main
  DW_AT_decl_file     : 1
  DW_AT_decl_line     : 7
  DW_AT_type          : <79>
  DW_AT_low_pc        : 0 0
  DW_AT_high_pc       : 0 0
  DW_AT_frame_base   : 1 byte block: 5e          (DW_OP_reg14)

<1><79>: Abbrev Number: 3 (DW_TAG_base_type)
  DW_AT_byte_size     : 2
  DW_AT_encoding      : 5          (signed)
  DW_AT_name          : int
```

This abbreviation is located at offset 0x62 from the start of the .debug_info section. This offset is in 'unpadded' bytes. To calculate the actual file offset, we can start with the file offset for the section (0xDC). The location of this entry is 0x62 unpadded bytes from the start of the section, or file offset 0xDC+(0x62*) == 0x1A0. Here is some of the raw, with padding, data from that area of the file:

```
00001a0  02 00 01 00 6d 00 61 00 69 00 6e 00 00 00 01 00
          002 \0 001 \0 m \0 a \0 i \0 n \0 \0 \0 001 \0
00001b0  07 00 79 00 00 00 00 00 00 00 00 00 00 00 00 00
          \a \0 y \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0
00001c0  00 00 00 00 00 00 00 00 00 00 01 00 5e 00 03 00
          \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 001 \0 ^ \0 003 \0
```

Without going to the specifics of the format, or knowing much about the sizes of each element, we can see roughly how it all lines up. Squint, to remove the padding, and you can find each of the values (high-lit above to help with the

squinting)!

Most integers are stored in SLEB128 or ULEB128, which stores values in a variable length. The value is chunked into 7-bit quantities and emitted in least endian form; if the remaining bits are all 0 we stop emitting the 7-bit chunks, if there is more data to follow we set the high-bit in the 8-bit value. Ie, the value 0x79 would be encoded in 1 byte (0x79) but the value 0x80 would be encoded in two bytes (0x80 0x01) - the first byte gives the low 7-bits (000.0000) and says 'more to come' and the next byte gives the next 7-bits (000.0001) and says 'finished'.

Note here that the value '0x79' which gives an offset for the function type (DW_TAG_subprogram) is unpadding.

To Make a Long Story Short

To decode the data from our DWARF sections when compiled for PIC24, dsPIC30F, or dsPIC33C/E/F devices; skip every other byte. Remember FILE offsets will refer to the padded offset, but DWARF offsets will refer to the unpadding offset.

Microchip MPLAB XC-DSC uses the same elf machine number for all dsPIC type devices. To determine if padding bytes require removal, examine the E_FLAGS field in the ELF header; bit 31 is clear to indicate that padding is included and set if it is not. All devices supported by MPLAB XC16 will have this bit clear.

The XC16 and XC-DSC compilers define the following flags in the ELF file header field:

```
#define EF_PIC30_PROC 0x0000FFFF /* 2 bytes in the e_flags field
                                for encoding processor ID */
#define EF_PIC30_NO_PHANTOM_BYTE 0x80000000 /* set if encoding has no phantom byte */
```

If an external tool is used to interpret DWARF debugging information created by the XC16 and/or XC-DSC compilers, the EF_PIC30_NO_PHANTOM_BYTE flag can be used to determine if phantom bytes in DWARF sections must be recognized and discarded.

For example:

- When EF_PIC30_NO_PHANTOM_BYTE is clear, DWARF sections include phantom bytes.
- When EF_PIC30_NO_PHANTOM_BYTE is set, DWARF sections do not include phantom bytes.

No labels